

ALY 6015 M3 Report - Thota, Sunil Raj.R

```
# Intermediate Analytics
# ALY 6015
# Module 3 - Regularization Assignment
# 02/03/2021
# Sunil Raj Thota
# NUID: 001099670

# Get and set the working directories
getwd()

## [1] "G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments"

setwd('G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments')
getwd()

## [1] "G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments"

# Installed the above packages into the work space
install.packages("datasets")
install.packages("plyr")
install.packages("dplyr")
install.packages("tidyr")
install.packages("ncvreg")
install.packages("biglasso")
install.packages("bigmemory")
install.packages("glmnet")

# Loaded the below libraries into the work space

library(plyr)
library(dplyr)
library(tidyr)
require(datasets)
library(biglasso)
library(bigmemory)
library(ncvreg)

data(mtcars)
attach(mtcars)
View(mtcars)
```

```
head(mtcars)
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

```
tail(mtcars)
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.7  0  1    5    2
## Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.9  1  1    5    2
## Ford Pantera L 15.8   8 351.0 264 4.22 3.170 14.5  0  1    5    4
## Ferrari Dino   19.7   6 145.0 175 3.62 2.770 15.5  0  1    5    6
## Maserati Bora  15.0   8 301.0 335 3.54 3.570 14.6  0  1    5    8
## Volvo 142E     21.4   4 121.0 109 4.11 2.780 18.6  1  1    4    2
```

```
str(mtcars)
```

```
## 'data.frame':   32 obs. of  11 variables:
## $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num  6 6 4 6 8 6 8 4 4 6 ...
## $ disp: num  160 160 108 258 360 ...
## $ hp : num  110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num  3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt : num  2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num  16.5 17 18.6 19.4 17 ...
## $ vs : num  0 0 1 1 0 1 0 1 1 1 ...
## $ am : num  1 1 1 0 0 0 0 0 0 0 ...
## $ gear: num  4 4 4 3 3 3 3 4 4 4 ...
## $ carb: num  4 4 1 1 2 1 4 2 2 4 ...
```

```
summary(mtcars)
```

```
##           mpg           cyl           disp           hp
## Min.   :10.40   Min.   :4.000   Min.   : 71.1   Min.   : 52.0
## 1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
## Median :19.20   Median :6.000   Median :196.3   Median :123.0
## Mean   :20.09   Mean   :6.188   Mean   :230.7   Mean   :146.7
## 3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
## Max.    :33.90   Max.    :8.000   Max.    :472.0   Max.    :335.0
##           drat           wt           qsec           vs
## Min.    :2.760   Min.    :1.513   Min.    :14.50   Min.    :0.0000
## 1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
## Median :3.695   Median :3.325   Median :17.71   Median :0.0000
## Mean    :3.597   Mean    :3.217   Mean    :17.85   Mean    :0.4375
## 3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
## Max.    :4.930   Max.    :5.424   Max.    :22.90   Max.    :1.0000
```

```
##           am           gear           carb
## Min.      :0.0000   Min.    :3.000   Min.     :1.000
## 1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
## Median :0.0000   Median :4.000   Median :2.000
## Mean     :0.4062   Mean    :3.688   Mean     :2.812
## 3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
## Max.     :1.0000   Max.    :5.000   Max.     :8.000
```

Let's perform some Regularization analysis and techniques using "mtcars" data set. This data set is readily available in the R Studio and can be loaded to the work space in R Studio. Or we can also install the packages by using install.packages("packagename") command. Once it is loaded we can use it in the code for further analysis and calculations.

Loaded the "mtcars" data into the work space. To reduce the repetitive usage of "mtcars" data set, "attach" is used to set it once throughout the work space. To View the diabetes Data set we use View() command, To observe the structure of the Data set we use str() command, and head () and tail() shows first and last few rows in the Data set. Summary() Provides the Descriptive Stats of the x variable in diabetes Data set.

```
y <- mtcars$hp
```

```
y
```

```
## [1] 110 110 93 110 175 105 245 62 95 123 123 180 180 180 205 215 230
## 66 52
```

```
## [20] 65 97 150 150 245 175 66 91 113 264 175 335 109
```

```
x <- data.matrix(mtcars[, c('mpg', 'wt', 'drat', 'qsec')])
```

```
x
```

```
##           mpg      wt drat  qsec
## Mazda RX4      21.0 2.620 3.90 16.46
## Mazda RX4 Wag  21.0 2.875 3.90 17.02
## Datsun 710      22.8 2.320 3.85 18.61
## Hornet 4 Drive  21.4 3.215 3.08 19.44
## Hornet Sportabout 18.7 3.440 3.15 17.02
## Valiant        18.1 3.460 2.76 20.22
## Duster 360     14.3 3.570 3.21 15.84
## Merc 240D      24.4 3.190 3.69 20.00
## Merc 230       22.8 3.150 3.92 22.90
## Merc 280       19.2 3.440 3.92 18.30
## Merc 280C      17.8 3.440 3.92 18.90
## Merc 450SE     16.4 4.070 3.07 17.40
## Merc 450SL     17.3 3.730 3.07 17.60
## Merc 450SLC    15.2 3.780 3.07 18.00
## Cadillac Fleetwood 10.4 5.250 2.93 17.98
## Lincoln Continental 10.4 5.424 3.00 17.82
## Chrysler Imperial 14.7 5.345 3.23 17.42
## Fiat 128       32.4 2.200 4.08 19.47
## Honda Civic    30.4 1.615 4.93 18.52
```

```
## Toyota Corolla      33.9 1.835 4.22 19.90
## Toyota Corona       21.5 2.465 3.70 20.01
## Dodge Challenger    15.5 3.520 2.76 16.87
## AMC Javelin         15.2 3.435 3.15 17.30
## Camaro Z28          13.3 3.840 3.73 15.41
## Pontiac Firebird    19.2 3.845 3.08 17.05
## Fiat X1-9           27.3 1.935 4.08 18.90
## Porsche 914-2       26.0 2.140 4.43 16.70
## Lotus Europa        30.4 1.513 3.77 16.90
## Ford Pantera L      15.8 3.170 4.22 14.50
## Ferrari Dino        19.7 2.770 3.62 15.50
## Maserati Bora       15.0 3.570 3.54 14.60
## Volvo 142E          21.4 2.780 4.11 18.60
```

```
linReglOLS <- lm(y ~ x)
linReglOLS
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Coefficients:
## (Intercept)      xmpg      xwt      xdrat      xqsec
##      473.779      -2.877      26.037      4.819     -20.751
```

```
summary(linReglOLS)
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -48.801 -16.007  -5.482  11.614  97.338
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  473.779    105.213   4.503 0.000116 ***
## xmpg         -2.877     2.381  -1.209 0.237319
## xwt          26.037    13.514   1.927 0.064600 .
## xdrat         4.819    15.952   0.302 0.764910
## xqsec        -20.751     3.993  -5.197 1.79e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 32.25 on 27 degrees of freedom
## Multiple R-squared:  0.8073, Adjusted R-squared:  0.7787
## F-statistic: 28.27 on 4 and 27 DF, p-value: 2.647e-09
```

Here "y" variable is taken as the response variable. Here "x" is assigned with a matrix of predictor variables

In this, we need to regress "y" on the predictors in "x" using Ordinary Least Squares(OLS). The regression model was taken between "y" and "x"

Summary() gives us the descriptive stats and hypothesis testing values Like Standard Error, p-Value, t-Value, r-squared value, f-Statistic, Degrees of Freedom, and etc.,

This model is used as a baseline model to collate with the next upcoming models

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.3
```

```
## Loaded glmnet 4.1
```

```
lambdaSeq <- 10 ^ seq(2, -2, by = -.1)
```

```
lambdaSeq
```

```
## [1] 100.00000000 79.43282347 63.09573445 50.11872336 39.81071706
## [6] 31.62277660 25.11886432 19.95262315 15.84893192 12.58925412
## [11] 10.00000000 7.94328235 6.30957344 5.01187234 3.98107171
## [16] 3.16227766 2.51188643 1.99526231 1.58489319 1.25892541
## [21] 1.00000000 0.79432823 0.63095734 0.50118723 0.39810717
## [26] 0.31622777 0.25118864 0.19952623 0.15848932 0.12589254
## [31] 0.10000000 0.07943282 0.06309573 0.05011872 0.03981072
## [36] 0.03162278 0.02511886 0.01995262 0.01584893 0.01258925
## [41] 0.01000000
```

```
ridgeFit <- glmnet(x, y, alpha = 0, lambda = lambdaSeq)
```

```
ridgeFit
```

```
##
```

```
## Call: glmnet(x = x, y = y, alpha = 0, lambda = lambdaSeq)
```

```
##
```

```
##      Df %Dev Lambda
## 1     4 64.58 100.000
## 2     4 68.02 79.430
## 3     4 70.91 63.100
## 4     4 73.28 50.120
## 5     4 75.18 39.810
## 6     4 76.65 31.620
## 7     4 77.78 25.120
## 8     4 78.62 19.950
## 9     4 79.24 15.850
## 10    4 79.68 12.590
## 11    4 80.00 10.000
## 12    4 80.23 7.943
## 13    4 80.38 6.310
## 14    4 80.49 5.012
```

```
## 15  4 80.57   3.981
## 16  4 80.62   3.162
## 17  4 80.65   2.512
## 18  4 80.68   1.995
## 19  4 80.69   1.585
## 20  4 80.70   1.259
## 21  4 80.71   1.000
## 22  4 80.72   0.794
## 23  4 80.72   0.631
## 24  4 80.72   0.501
## 25  4 80.72   0.398
## 26  4 80.72   0.316
## 27  4 80.72   0.251
## 28  4 80.72   0.200
## 29  4 80.73   0.158
## 30  4 80.73   0.126
## 31  4 80.73   0.100
## 32  4 80.73   0.079
## 33  4 80.73   0.063
## 34  4 80.73   0.050
## 35  4 80.73   0.040
## 36  4 80.73   0.032
## 37  4 80.73   0.025
## 38  4 80.73   0.020
## 39  4 80.73   0.016
## 40  4 80.73   0.013
## 41  4 80.73   0.010
```

```
summary(ridgeFit)
```

```
##           Length Class      Mode
## a0           41    -none-  numeric
## beta        164 dgCMatrix S4
## df           41    -none-  numeric
## dim           2    -none-  numeric
## lambda        41    -none-  numeric
## dev.ratio     41    -none-  numeric
## nulldev        1    -none-  numeric
## npasses        1    -none-  numeric
## jerr           1    -none-  numeric
## offset         1    -none-  logical
## call           5    -none-   call
## nobs           1    -none-  numeric
```


Setting the range of Lambda values and Using glmnet() method to build the ridge regression in R. Checking the model using the Summary()

```
modelLASSO <- glmnet(x, y, alpha = 1)
modelLASSO
```

```
##
## Call:  glmnet(x = x, y = y, alpha = 1)
##
##      Df  %Dev Lambda
## 1    0  0.00 52.380
## 2    1 10.23 47.730
## 3    2 19.52 43.490
## 4    2 29.45 39.620
## 5    2 37.71 36.100
## 6    2 44.56 32.900
## 7    2 50.24 29.970
## 8    2 54.97 27.310
## 9    2 58.89 24.880
## 10   2 62.14 22.670
## 11   2 64.84 20.660
## 12   2 67.09 18.820
## 13   3 69.32 17.150
## 14   3 71.25 15.630
## 15   3 72.84 14.240
## 16   3 74.17 12.970
## 17   3 75.27 11.820
## 18   3 76.19 10.770
## 19   3 76.95  9.815
## 20   3 77.58  8.943
## 21   3 78.10  8.148
## 22   3 78.53  7.424
## 23   3 78.90  6.765
## 24   3 79.19  6.164
## 25   3 79.44  5.616
## 26   3 79.65  5.117
## 27   3 79.82  4.663
## 28   3 79.96  4.249
## 29   3 80.08  3.871
## 30   3 80.18  3.527
## 31   3 80.26  3.214
## 32   3 80.33  2.928
## 33   3 80.39  2.668
## 34   3 80.43  2.431
## 35   3 80.47  2.215
## 36   3 80.50  2.018
## 37   3 80.53  1.839
## 38   3 80.55  1.676
## 39   3 80.57  1.527
## 40   3 80.59  1.391
## 41   3 80.60  1.268
## 42   3 80.61  1.155
## 43   3 80.62  1.052
## 44   3 80.62  0.959
## 45   3 80.63  0.874
## 46   3 80.64  0.796
```

```
## 47 3 80.64 0.725
## 48 4 80.65 0.661
## 49 4 80.67 0.602
## 50 4 80.68 0.549
## 51 4 80.68 0.500
## 52 4 80.69 0.456
## 53 4 80.70 0.415
## 54 4 80.70 0.378
## 55 4 80.71 0.345
## 56 4 80.71 0.314
## 57 4 80.71 0.286
## 58 4 80.71 0.261
## 59 4 80.72 0.238
## 60 4 80.72 0.216
## 61 4 80.72 0.197
## 62 4 80.72 0.180
## 63 4 80.72 0.164
## 64 4 80.72 0.149
```

```
plot(modelLASSO,
      xvar = "norm",
      label = TRUE)
```



LASSO regression is performed and for that to happen we use "glmnet" package from the packages tab to install or simply use `install.packages("glmnet")` command

Now, let's load the "glmnet" in our work space to regularize the model using LASSO and plot it using plot(). This plot indicates at which stage each coefficients shrinks to 0. and the lines depicts the values used by various other coefficients


```
cvModel <- cv.glmnet(x, y, alpha = 1)
cvModel

##
## Call:  cv.glmnet(x = x, y = y, alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min   2.431    34   1250 362.4      3
## 1se  14.239    15   1573 485.3      3

bestLambda <- cvModel$lambda.min
bestLambda

## [1] 2.431182

plot(cvModel)
```



Here, Cross Validation is used to get the best value of lambda and plot the curve using plot(). It is possible with cv.glmnet() method. nlambdas signifies the number of lambda values in sequence. In general, nlambdas values must be above 100.

Let's find optimal lambda value that minimizes test MSE and perform K-Fold Cross validation to find optimal lambda value. At last, let's produce the plot of test MSE by lambda value.

From the plot we can depict that the value of lambda increased when the number of selected variables narrows down. This tells that higher the value of lambda, more shrink the selection is. Now, we find the min. value of lambda to get the best fit

```
lambdaWithOneSE <- cvModel$lambda.1se  
lambdaWithOneSE
```

```
## [1] 14.23948
```

```
latestFit <- glmnet(  
  x = x,  
  y = y,  
  alpha = 1,  
  lambda = lambdaWithOneSE  
)
```

```
latestFit$beta
```

```
## 4 x 1 sparse Matrix of class "dgCMatrix"  
##           s0  
## mpg    -4.024702  
## wt      5.766737  
## drat     .  
## qsec  -12.842480
```

Here, we use the minimum lambda value again in glmnet() function to get the best latest fit. Now we use a higher value of lambda that is within one standard error of the minimum to check its effect on shrinkage.

There are 1 coefficients namely "drat" whose values have become 0. It's clear that this variable is not so necessary to determine the value of "y". LASSO tells that only 3 variables are necessary on which y depends. Thus the shrinkage increases.

```
bestModel <- glmnet(x, y, alpha = 1, lambda = bestLambda)  
coef(bestModel)
```

```
## 5 x 1 sparse Matrix of class "dgCMatrix"  
##           s0  
## (Intercept) 485.152675  
## mpg        -2.936266  
## wt         21.698919  
## drat         .  
## qsec       -19.569135
```

```
newObs <- matrix(c(21, 2.1, 3.6, 18.0), nrow = 1, ncol = 4)
newObs
```

```
##      [,1] [,2] [,3] [,4]
## [1,]   21  2.1  3.6   18
```

```
predict(bestModel, s = bestLambda, newx = newObs)
```

```
##           1
## [1,] 116.8144
```

```
yPred <- predict(bestModel, s = bestLambda, newx = x)
```

```
sstValue <- sum((y - mean(y)) ^ 2)
sseValue <- sum((yPred - y) ^ 2)
```

```
rSquaredVal <- 1 - sseValue / sstValue
rSquaredVal
```

```
## [1] 0.8043193
```

To find the coefficients of best model, let's define a new observation and use LASSO regression model to predict response value. Use fitted best model to make predictions. Let's find SST, SSE, and R-Squared values for the new observation