

Thota, Sunil Raj - LASSO Regression in R Practice.R

```
# Intermediate Analytics
# ALY 6015
# Module 3 - LASSO Regression in R Practice
# 02/03/2021
# Sunil Raj Thota
# NUID: 001099670

# Get and set the working directories
getwd()

## [1] "G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments"

setwd('G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments')
getwd()

## [1] "G:/NEU/Coursework/2021 Q1 Winter/ALY 6015 IA/Discussions &
Assignments"

# Installed the above packages into the work space

install.packages("datasets")
install.packages("plyr")
install.packages("dplyr")
install.packages("tidyr")
install.packages("ncvreg")
install.packages("biglasso")
install.packages("bigmemory")
install.packages("glmnet")
install.packages("lars")

# Loaded the below libraries into the work space

library(plyr)
library(dplyr)
library(tidyr)
require(datasets)
library(biglasso)
library(bigmemory)
library(ncvreg)
```

Exercise 1

```
library(lars)

data(diabetes)
attach(diabetes)

View(diabetes)
str(diabetes)

## 'data.frame':    442 obs. of  3 variables:
## $ x : 'AsIs' num [1:442, 1:10] 0.03808 -0.00188 0.0853 -0.08906 0.00538
## ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:10] "age" "sex" "bmi" "map" ...
## $ y : num 151 75 141 206 135 97 138 63 110 310 ...
## $ x2: 'AsIs' num [1:442, 1:64] 0.03808 -0.00188 0.0853 -0.08906 0.00538
## ...
## .. attr(*, ".Names")= chr [1:28288] "age" "age" "age" "age" ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:442] "1" "2" "3" "4" ...
## .. ..$ : chr [1:64] "age" "sex" "bmi" "map" ...

head(diabetes)

##           x.age           x.sex           x.bmi           x.map           x.tc
x.ldl
## 1  0.038075906  0.050680119  0.061696207  0.021872355 -0.044223498 -
0.034820763
## 2 -0.001882017 -0.044641637 -0.051474061 -0.026327835 -0.008448724 -
0.019163340
## 3  0.085298906  0.050680119  0.044451213 -0.005670611 -0.045599451 -
0.034194466
## 4 -0.089062939 -0.044641637 -0.011595015 -0.036656447  0.012190569
0.024990593
## 5  0.005383060 -0.044641637 -0.036384692  0.021872355  0.003934852
0.015596140
## 6 -0.092695478 -0.044641637 -0.040695940 -0.019442093 -0.068990650 -
0.079287844
##           x.hdl           x.tch           x.ltg           x.glu    y           x2.age
## 1 -0.043400846 -0.002592262  0.019908421 -0.017646125 151  0.0380759064
## 2  0.074411564 -0.039493383 -0.068329744 -0.092204050  75 -0.0018820165
## 3 -0.032355932 -0.002592262  0.002863771 -0.025930339 141  0.0852989063
## 4 -0.036037570  0.034308859  0.022692023 -0.009361911 206 -0.0890629394
## 5  0.008142084 -0.002592262 -0.031991445 -0.046640874 135  0.0053830604
## 6  0.041276824 -0.076394504 -0.041180385 -0.096346157  97 -0.0926954778
##           x2.sex           x2.bmi           x2.map           x2.tc           x2.ldl
## 1  0.0506801187  0.0616962065  0.0218723550 -0.0442234984 -0.0348207628
## 2 -0.0446416365 -0.0514740612 -0.0263278347 -0.0084487241 -0.0191633397
## 3  0.0506801187  0.0444512133 -0.0056706106 -0.0455994513 -0.0341944659
```

```

## 4 -0.0446416365 -0.0115950145 -0.0366564468 0.0121905688 0.0249905934
## 5 -0.0446416365 -0.0363846922 0.0218723550 0.0039348516 0.0155961395
## 6 -0.0446416365 -0.0406959405 -0.0194420933 -0.0689906499 -0.0792878444
##      x2.hdl      x2.tch      x2.ltg      x2.glu      x2.age^2
## 1 -0.0434008457 -0.0025922620 0.0199084209 -0.0176461252 -0.0148551625
## 2 0.0744115641 -0.0394933829 -0.0683297436 -0.0922040496 -0.0412915429
## 3 -0.0323559322 -0.0025922620 0.0028637705 -0.0259303390 0.0916434391
## 4 -0.0360375700 0.0343088589 0.0226920226 -0.0093619113 0.1036403301
## 5 0.0081420836 -0.0025922620 -0.0319914449 -0.0466408736 -0.0408265979
## 6 0.0412768238 -0.0763945038 -0.0411803852 -0.0963461565 0.1157092549
##      x2.bmi^2      x2.map^2      x2.tc^2      x2.ldl^2      x2.hdl^2
## 1 0.0225045739 -0.0310446765 -0.0043311197 -0.0137399243 -0.0046314248
## 2 0.0056427733 -0.0273076609 -0.0309389016 -0.0248010319 0.0400365241
## 3 -0.0041764214 -0.0388099038 -0.0025859366 -0.0143055611 -0.0148614560
## 4 -0.0310170859 -0.0159874156 -0.0298483890 -0.0214340114 -0.0117828858
## 5 -0.0136807215 -0.0310446765 -0.0317282077 -0.0264236384 -0.0268506700
## 6 -0.0088370145 -0.0327918526 0.0352626533 0.0526602897 -0.0068304035
##      x2.tch^2      x2.ltg^2      x2.glu^2      x2.age:sex      x2.age:bmi
## 1 -0.0304484629 -0.0288162192 -0.0275255618 0.0328649758 0.0405716741
## 2 -0.0094854824 0.0371612444 0.0880219609 -0.0066099928 -0.0067648038
## 3 -0.0304484629 -0.0348099250 -0.0224326123 0.0840517453 0.0708891293
## 4 -0.0146503349 -0.0269850701 -0.0306820952 0.0766292449 0.0129034054
## 5 -0.0304484629 -0.0191324512 -0.0012284185 -0.0135465959 -0.0129173209
## 6 0.0482386067 -0.0087497144 0.0990402558 0.0800975464 0.0704832806
##      x2.age:map      x2.age:tc      x2.age:ldl      x2.age:hdl      x2.age:tch
## 1 0.0016606410 -0.0465532511 -0.0382447104 -0.0345115069 -0.0121122609
## 2 -0.0159342243 -0.0117287997 -0.0096555406 0.0006995477 -0.0083689963
## 3 -0.0279128687 -0.0917442739 -0.0716415163 -0.0602920920 -0.0147605279
## 4 0.0562904032 -0.0342989420 -0.0571356258 0.0786805432 -0.0760817036
## 5 -0.0144024121 -0.0116206046 -0.0086502423 0.0049801694 -0.0102788452
## 6 0.0234365273 0.1189685396 0.1438718189 -0.0851146968 0.1432199339
##      x2.age:ltg      x2.age:glu      x2.sex:bmi      x2.sex:map      x2.sex:tc
## 1 0.0030648916 -0.0302775066 0.0621030123 0.0122820371 -0.0485722318
## 2 -0.0102016795 -0.0113801440 0.0445181909 0.0137392710 0.0062226212
## 3 -0.0077635174 -0.0646990887 0.0435615292 -0.0181579597 -0.0500315236
## 4 -0.0555091413 0.0033785934 0.0067497817 0.0237941848 -0.0130586592
## 5 -0.0165418441 -0.0208710562 0.0302274415 -0.0331836602 -0.0053461470
## 6 0.0675437504 0.1843686827 0.0343105127 0.0070359951 0.0627810439
##      x2.sex:ldl      x2.sex:hdl      x2.sex:tch      x2.sex:ltg      x2.sex:glu
## 1 -0.0441240238 -0.0308042981 -0.0195479919 0.0142268228 -0.0293859648
## 2 0.0112617659 -0.0565675488 0.0224021267 0.0575880019 0.0784642525
## 3 -0.0434530875 -0.0179545693 -0.0195479919 -0.0041216840 -0.0384231554
## 4 -0.0304033769 0.0566194244 -0.0505546014 -0.0287218500 -0.0011399371
## 5 -0.0215384529 0.0113446351 -0.0140762373 0.0231308241 0.0346819482
## 6 0.0679972794 -0.0226114568 0.0588804908 0.0318440815 0.0824444620
##      x2.bmi:map      x2.bmi:tc      x2.bmi:ldl      x2.bmi:hdl      x2.bmi:tch
## 1 0.0090008988 -0.0717180778 -0.0603176794 -0.0413575386 -0.0240342004
## 2 0.0091148702 -0.0028355367 0.0087097314 -0.0671553344 0.0240456899
## 3 -0.0226918084 -0.0564432291 -0.0464818380 -0.0136167525 -0.0230540270
## 4 -0.0092925186 -0.0153834201 -0.0193921008 0.0279274131 -0.0292499537

```

```
## 5 -0.0334523099 -0.0154230197 -0.0255070025 0.0119441348 -0.0184594644
## 6 -0.0020460266 0.0488321410 0.0580412284 -0.0190229456 0.0476394964
##      x2.bmi:ltg      x2.bmi:glu      x2.map:tc      x2.map:ldl      x2.map:hd1
## 1 0.0048261936 -0.0410278367 -0.0327209536 -0.0257474212 -0.0112074291
## 2 0.0552994440 0.0806093115 -0.0070399803 0.0018462404 -0.0319794277
## 3 -0.0194513972 -0.0423606977 -0.0062598621 -0.0049233843 0.0120934525
## 4 -0.0280603703 -0.0160690142 -0.0214874364 -0.0291135216 0.0354925517
## 5 0.0034088636 0.0170453175 -0.0099837004 -0.0017149265 0.0119825521
## 6 0.0146962189 0.0634061398 0.0171122204 0.0244459437 -0.0081883427
##      x2.map:tch      x2.map:ltg      x2.map:glu      x2.tc:ldl      x2.tc:hd1
## 1 -0.0138251131 -0.0101938361 -0.0257784633 -0.0068689647 0.0398230899
## 2 0.0098745200 0.0203696547 0.0313619644 -0.0262353123 -0.0164622948
## 3 -0.0122818754 -0.0203183065 -0.0149534838 -0.0065969775 0.0300168640
## 4 -0.0397827564 -0.0385992765 -0.0109701272 -0.0242291837 -0.0122792654
## 5 -0.0138251131 -0.0356386917 -0.0386583542 -0.0276482686 -0.0018670731
## 6 0.0195036026 -0.0020081382 0.0201031994 0.0483666130 -0.0654803859
##      x2.tc:tch      x2.tc:ltg      x2.tc:glu      x2.ldl:hd1      x2.ldl:tch
## 1 -0.0193030537 -0.0428915072 0.0009629700 0.0423549304 -0.0220378305
## 2 -0.0155012002 -0.0123430995 0.0009326831 -0.0212564243 -0.0115642516
## 3 -0.0192411418 -0.0271777643 0.0098716066 0.0335869646 -0.0220633407
## 4 -0.0140331587 -0.0186440420 -0.0188580962 -0.0098784247 -0.0099839514
## 5 -0.0214699727 -0.0270791696 -0.0203958711 0.0123759491 -0.0240914053
## 6 0.0701909475 0.0350969957 0.1309600895 -0.0612519836 0.0717191451
##      x2.ldl:ltg      x2.ldl:glu      x2.hdl:tch      x2.hdl:ltg      x2.hdl:glu
## 1 -0.0311245646 -0.0009221095 0.0334936252 0.0008521487 0.0311502576
## 2 0.0129733755 0.0237834359 -0.0238146613 -0.0945055990 -0.1403775894
## 3 -0.0180161691 0.0049134553 0.0329558784 0.0182807986 0.0327952439
## 4 -0.0033727555 -0.0191092832 0.0081586184 0.0018977323 0.0215138966
## 5 -0.0268464903 -0.0296874121 0.0309841399 0.0144891320 0.0053856590
## 6 0.0560369456 0.1496630223 -0.0278444433 -0.0180308997 -0.0755127518
##      x2.tch:ltg      x2.tch:glu      x2.ltg:glu
## 1 -0.0281911757 -0.0176581553 -0.0277936831
## 2 0.0252977155 0.0530335390 0.1040132768
## 3 -0.0273318271 -0.0172359590 -0.0223037368
## 4 -0.0120454909 -0.0248722040 -0.0250419367
## 5 -0.0255745144 -0.0161804685 0.0087351987
## 6 0.0339989574 0.1261465729 0.0577886517
```

Let's perform some Regularization analysis and techniques using "diabetes" data set. To get this data set we need to install the 'lars' package from the packages tab which is right side to the work space in R Studio.

Or we can also install the packages by using `install.packages("package name")` command. Once it is loaded we can use it in the code for further analysis and calculations.

Loaded the "lars" library into the work space. Loaded the diabetes Data set into the Environment. To reduce the repetitive usage of "x" in "diabetes" data set, "attach" is used to set it once throughout the work space.

To View the diabetes Data set we use View() command, To observe the structure of the Data set we use str() command, and head () and tail() shows first and last few rows in the Data set

Exercise 2

```
summary(x)
```







```
##      age              sex              bmi              map
## Min.   :-0.107226   Min.    :-0.04464   Min.     :-0.090275   Min.     :-
0.112400
## 1st Qu.: -0.037299   1st Qu.: -0.04464   1st Qu.: -0.034229   1st Qu.: -
0.036656
## Median : 0.005383   Median : -0.04464   Median : -0.007284   Median : -
0.005671
## Mean   : 0.000000   Mean    : 0.00000   Mean     : 0.000000   Mean      :
0.000000
## 3rd Qu.: 0.038076   3rd Qu.: 0.05068   3rd Qu.: 0.031248   3rd Qu.:
0.035644
## Max.    : 0.110727   Max.     : 0.05068   Max.     : 0.170555   Max.      :
0.132044
##      tc              ldl              hdl
## Min.   :-0.126781   Min.    :-0.115613   Min.     :-0.102307
## 1st Qu.: -0.034248   1st Qu.: -0.030358   1st Qu.: -0.035117
## Median : -0.004321   Median : -0.003819   Median : -0.006584
## Mean   : 0.000000   Mean     : 0.000000   Mean     : 0.000000
## 3rd Qu.: 0.028358   3rd Qu.: 0.029844   3rd Qu.: 0.029312
## Max.    : 0.153914   Max.     : 0.198788   Max.     : 0.181179
##      tch              ltg              glu
## Min.   :-0.076395   Min.    :-0.126097   Min.     :-0.137767
## 1st Qu.: -0.039493   1st Qu.: -0.033249   1st Qu.: -0.033179
## Median : -0.002592   Median : -0.001948   Median : -0.001078
## Mean   : 0.000000   Mean     : 0.000000   Mean     : 0.000000
## 3rd Qu.: 0.034309   3rd Qu.: 0.032433   3rd Qu.: 0.027917
## Max.    : 0.185234   Max.     : 0.133599   Max.     : 0.135612
```

```
par(mfcol = c(2, 3))
for (idx in 1:10)
{
  if (idx == 1) {
    xLabel = "ages"
  }
  if (idx == 2) {
    xLabel = "sex"
  }
  if (idx == 3) {
    xLabel = "bmi"
  }
  if (idx == 4) {
    xLabel = "map"
  }
}
```

```

}
if (idx == 5) {
  xlab = "tc"
}
if (idx == 6) {
  xLabel = "ldl"
}
if (idx == 7) {
  xLabel = "hdl"
}
if (idx == 8) {
  xLabel = "tch"
}
if (idx == 9) {
  xLabel = "ltg"
}
if (idx == 10) {
  xLabel = "glu"
}
plot(x[, idx], y, xlab = xLabel)
abline(lm(y ~ x[, idx]), col = "tomato")
}

```




Summary() Provides the Descriptive Stats of the x variable in diabetes Data set. Once this is done, let's find out the list of elements which are present in "x" using looping concept (for loop).

We noticed 10 variables from the statistics given in the summary. Now, let's plot Scatter plots for these variables and stack them in 2 rows and 5 columns

To generate separate Scatter Plots with the line of best fit for all the predictors in "x" as horizontal axis vs "y" as vertical axis.

In this for loop I used 'idx' as index value for "x" variable and increased the value of 'idx' by iterating the loop 10 times.

I have included conditional statements to allocate individual x-labels to their respective plots. abline() is used to plot a linear regression line on these plots.



Exercise 3

```
linReg1OLS <- lm(y ~ x)
linReg1OLS
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ x)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      xage      xsex      xbm1      xmap
```

```
xtc
```

```
##      152.13      -10.01     -239.82      519.84      324.39      -
```

```
792.18
##          xldl          xhdl          xtch          xltg          xglu
##          476.75          101.04          177.06          751.28          67.63
```

```
summary(linReg1OLS)
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -155.829  -38.534   -0.227   37.806  151.355
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  152.133      2.576   59.061 < 2e-16 ***
## xage         -10.012      59.749   -0.168  0.867000
## xsex         -239.819     61.222  -3.917  0.000104 ***
## xbmi          519.840     66.534   7.813  4.30e-14 ***
## xmap          324.390     65.422   4.958  1.02e-06 ***
## xtc          -792.184    416.684  -1.901  0.057947 .
## xldl          476.746    339.035   1.406  0.160389
## xhdl          101.045    212.533   0.475  0.634721
## xtch          177.064    161.476   1.097  0.273456
## xltg          751.279    171.902   4.370  1.56e-05 ***
## xglu           67.625     65.984   1.025  0.305998
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 54.15 on 431 degrees of freedom
## Multiple R-squared:  0.5177, Adjusted R-squared:  0.5066
## F-statistic: 46.27 on 10 and 431 DF, p-value: < 2.2e-16
```

In this, we need to regress "y" on the predictors in "x" using Ordinary Least Squares(OLS). The regression model was taken between "y" and "x"

Summary() gives us the descriptive stats and hypothesis testing values like Standard Error, p-Value, t-Value, r-squared value, f-Statistic, Degrees of Freedom, and etc.,

This model is used as a baseline model to collate with the next upcoming models

Exercise 4

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.3
```


```
## Loaded glmnet 4.1
```



```

modellASSO <- glmnet(x, y, alpha = 1)
plot(modellASSO,
      xvar = "norm",
      label = TRUE)

```



LASSO regression is performed and for that to happen we use "glmnet" package from the packages tab to install or simply use `install.packages("glmnet")` command


Now, Let's load the "glmnet" in our work space to regularize the model using LASSO and plot it using `plot()`. This plot indicates at which stage each coefficients shrinks to 0. and the lines depicts the values used by various other coefficients

Exercise 5

```

crossValidFit <- cv.glmnet(
  x = x,
  y = y,
  alpha = 1,
  nlambda = 1000
)
plot(crossValidFit)

```



```
minLambda <- crossValidFit$lambda.min
minLambda
```

```
## [1] 0.8730528
```

Here, Cross Validation is used to get the best value of lambda and plot the curve using plot(). It is possible with cv.glmnet() method. nlambda signifies the number of lambda values in sequence. In general, nlambda values must be above 100.

From the plot we can depict that the value of lambda increased when the number of selected variables narrows down. This tells that higher the value of lambda, more shrink the selection is. Now, we find the min. value of lambda to get the best fit

Exercise 6

```
estBetaMatFit <- glmnet(
  x = x,
  y = y,
  alpha = 1,
  lambda = minLambda
)
estBetaMatFit$beta
```

```
## 10 x 1 sparse Matrix of class "dgCMatrix"
##                s0
```

```

## age      .
## sex -201.065217
## bmi  522.602018
## map  299.095320
## tc   -111.562032
## ldl    .
## hdl -219.228019
## tch    8.497895
## ltg  516.016160
## glu   55.828353

estBetaMatFit$lambda

## [1] 0.8730528

# Here, we use the minimum lambda value again in glmnet() function to get the best fit

# There are 3 coefficients namely age, ldl, tch whose values have become 0. It's clear that these variables are not so necessary to determine the value of "y".

# Exercise 7

# Now we use a higher value of lambda that is within one standard error of the minimum to check its effect on shrinkage.

lambdaWithOneSE <- crossValidFit$lambda.1se
lambdaWithOneSE

## [1] 6.885531

latestFit <- glmnet(
  x = x,
  y = y,
  alpha = 1,
  lambda = lambdaWithOneSE
)

latestFit$beta

## 10 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## age      .
## sex      .
## bmi  500.1103
## map  182.4389
## tc      .
## ldl      .
## hdl -105.0984
## tch      .

```

```
## ltg 434.5747
## glu .
```

Here, we use the minimum lambda value again in glmnet() function to get the best latest fit

There are 6 coefficients namely age, sx, tc, ldl, tch, and glu whose values have become 0. It's clear that these variables are not so necessary to determine the value of "y". LASSO tells that only 4 variables are necessary on which y depends. Thus the shrinkage increases

Exercise 8

```
linReg1OLS2 <- lm(y ~ x2)
summary(linReg1OLS2)
```

```
##
## Call:
## lm(formula = y ~ x2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -158.216  -30.809   -3.857   31.348  153.946
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   152.133     2.532   60.086 < 2e-16 ***
## x2age         50.721     65.513    0.774  0.4393
## x2sex        -267.344     65.270   -4.096 5.15e-05 ***
## x2bmi         460.721     84.601    5.446 9.32e-08 ***
## x2map         342.933     72.447    4.734 3.13e-06 ***
## x2tc        -3599.542    60575.187  -0.059  0.9526
## x2ldl         3028.281    53238.699   0.057  0.9547
## x2hdl         1103.047    22636.179   0.049  0.9612
## x2tch          74.937     275.807   0.272  0.7860
## x2ltg         1828.210    19914.504   0.092  0.9269
## x2glu          62.754     70.398   0.891  0.3733
## x2age^2        67.691     69.470   0.974  0.3305
## x2bmi^2        45.849     83.288   0.550  0.5823
## x2map^2        -8.460     71.652  -0.118  0.9061
## x2tc^2        6668.449    7059.159   0.945  0.3454
## x2ldl^2       3583.174    5326.148   0.673  0.5015
## x2hdl^2       1731.821    1590.574   1.089  0.2769
## x2tch^2        773.374     606.967   1.274  0.2034
## x2ltg^2       1451.581    1730.103   0.839  0.4020
## x2glu^2        114.149     94.122   1.213  0.2260
## x2age:sex      148.678     73.407   2.025  0.0435 *
## x2age:bmi     -18.052     79.620  -0.227  0.8208
## x2age:map       18.534     76.303   0.243  0.8082
## x2age:tc      -158.891    617.109  -0.257  0.7970
```

```

## x2age:ldl      -67.285      494.527    -0.136     0.8918
## x2age:hdl      209.245      280.614     0.746     0.4563
## x2age:tch      184.960      210.330     0.879     0.3798
## x2age:ltg      124.667      223.765     0.557     0.5778
## x2age:glu       62.575       80.377     0.779     0.4367
## x2sex:bmi       64.612       77.902     0.829     0.4074
## x2sex:map       88.472       74.744     1.184     0.2373
## x2sex:tc       433.598      590.709     0.734     0.4634
## x2sex:ldl     -352.823      468.951    -0.752     0.4523
## x2sex:hdl     -124.731      273.870    -0.455     0.6491
## x2sex:tch     -131.223      199.714    -0.657     0.5115
## x2sex:ltg     -118.995      226.493    -0.525     0.5996
## x2sex:glu       45.758       73.650     0.621     0.5348
## x2bmi:map      154.720       86.340     1.792     0.0739 .
## x2bmi:tc     -302.045      667.930    -0.452     0.6514
## x2bmi:ldl      241.540      561.026     0.431     0.6671
## x2bmi:hdl      121.942      329.884     0.370     0.7118
## x2bmi:tch     -33.445      230.836    -0.145     0.8849
## x2bmi:ltg      114.673      255.987     0.448     0.6544
## x2bmi:glu       23.377       91.037     0.257     0.7975
## x2map:tc       478.303      682.264     0.701     0.4837
## x2map:ldl     -326.740      574.317    -0.569     0.5697
## x2map:hdl     -187.305      309.589    -0.605     0.5455
## x2map:tch     -58.294      198.601    -0.294     0.7693
## x2map:ltg     -154.795      271.966    -0.569     0.5696
## x2map:glu     -133.476       91.314    -1.462     0.1447
## x2tc:ldl     -9313.775     11771.220    -0.791     0.4293
## x2tc:hdl     -3932.025      3816.572    -1.030     0.3036
## x2tc:tch     -2205.910      1761.843    -1.252     0.2113
## x2tc:ltg     -3801.442     13166.091    -0.289     0.7729
## x2tc:glu     -176.295       595.459    -0.296     0.7673
## x2ldl:hdl      2642.645      3165.926     0.835     0.4044
## x2ldl:tch      1206.822      1470.512     0.821     0.4123
## x2ldl:ltg      2773.697     10960.214     0.253     0.8004
## x2ldl:glu       85.626       505.102     0.170     0.8655
## x2hdl:tch      1188.406      1002.242     1.186     0.2365
## x2hdl:ltg      1467.845      4609.793     0.318     0.7503
## x2hdl:glu      217.541       296.749     0.733     0.4640
## x2tch:ltg      389.805       624.671     0.624     0.5330
## x2tch:glu      235.693       235.064     1.003     0.3167
## x2ltg:glu       83.525       264.726     0.316     0.7525
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 53.23 on 377 degrees of freedom
## Multiple R-squared:  0.5924, Adjusted R-squared:  0.5233
## F-statistic: 8.563 on 64 and 377 DF, p-value: < 2.2e-16


```

In this, we need to regress "y" on the predictors in "x2" using Ordinary Least Squares(OLS). The regression model was taken between "y" and "x"2

Summary() gives us the descriptive stats and hypothesis testing values like Standard Error, p-Value, t-Value, r-squared value, f-Statistic, Degrees of Freedom, and etc., From this we can see that there are more number of predictors in x2 than x.

Exercise 9

```
modelLASSO2 <- glmnet(x2, y, alpha = 1)
plot(modelLASSO2, xvar = "norm", label = TRUE)
```




LASSO regression is performed and for that to happen we use "glmnet" package from the packages tab to install or simply use install.packages("glmnet") command

Now, Let's Load the "glmnet" in our work space to regularize the model using LASSO and plot it using plot(). This plot are complex and the lines depicts the values used by various other coefficients

Exercise 10

```
crossValidFit2 <- cv.glmnet(
  x = x2,
  y = y,
  alpha = 1,
  nlambda = 1000
)
```

```
plot(crossValidFit2)
```



```
minLambda2 <- crossValidFit2$lambda.min

estBetaMatFit2 <- glmnet(
  x = x2,
  y = y,
  alpha = 1,
  lambda = minLambda2
)
estBetaMatFit2$beta

## 64 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## age      .
## sex    -116.457385
## bmi     501.478323
## map     254.146483
## tc       .
## ldl      .
## hdl    -190.527089
## tch      .
## ltg     468.227778
## glu     19.670815
## age^2    9.718022
## bmi^2   39.571614
```

```
## map^2      .
## tc^2       .
## ldl^2      .
## hdl^2      .
## tch^2      .
## ltg^2      .
## glu^2      71.215436
## age:sex    109.359288
## age:bmi    .
## age:map    30.185352
## age:tc     .
## age:ldl    .
## age:hdl    .
## age:tch    .
## age:ltg    9.451642
## age:glu    11.123118
## sex:bmi    .
## sex:map    1.629466
## sex:tc     .
## sex:ldl    .
## sex:hdl    .
## sex:tch    .
## sex:ltg    .
## sex:glu    .
## bmi:map    86.745375
## bmi:tc     .
## bmi:ldl    .
## bmi:hdl    .
## bmi:tch    .
## bmi:ltg    .
## bmi:glu    .
## map:tc     .
## map:ldl    .
## map:hdl    .
## map:tch    .
## map:ltg    .
## map:glu    .
## tc:ldl     .
## tc:hdl     .
## tc:tch     .
## tc:ltg     .
## tc:glu     .
## ldl:hdl    .
## ldl:tch    .
## ldl:ltg    .
## ldl:glu    .
## hdl:tch    .
## hdl:ltg    .
## hdl:glu    .
## tch:ltg    .
```



```
## tch:glu      .  
## ltg:glu      .
```

```
estBetaMatFit2$lambda
```

```
## [1] 2.921225
```

Here, Cross Validation is used to get the best value of lambda and plot the curve using plot(). It is possible with cv.glmnet() method. nlambda signifies the number of lambda values in sequence. In general, n lambda values must be above 100.

From the plot we can depict that the value of lambda increased when the number of selected variables narrows down. This tells that higher the value of lambda, more shrink the selection is. Now, we find the min. value of lambda to get the best fit

Here, we use the minimum lambda value again in glmnet() function to get the best fit. There are 50 coefficients whose values have become 0. It's clear that these variables are not so necessary to determine the value of "y". With this it shrinkage's the variables and it regularizes the model.